

Biomedical Named Entity Recognition for Clinical Notes: A Comparative Study

Kazi Alif Haider

March 2026

Contents

1	Introduction	3
2	SciSpacy Pipeline	3
3	Transformer-Based NER Pipeline	3
4	Implementation Challenges	4
5	Experimental Environment	4
6	Performance Comparison	5
6.1	Runtime Statistics	5
6.2	NER Model Comparison	6
7	Summary of Findings	6

Abstract

Biomedical Named Entity Recognition (NER) is an important task in clinical natural language processing, enabling the extraction of structured medical concepts from unstructured clinical text. This study presents a practical comparison of multiple biomedical NER approaches applied to clinical notes. Three different methods were evaluated: a SciSpacy-based pipeline and transformer-based models built on BioBERT/PubMedBERT and ClinicalBERT architectures. Experiments were conducted on approximately 5,000 clinical notes collected from publicly available medical transcription samples. The comparison focuses on practical implementation aspects including handling of long clinical notes, tokenization behavior, engineering complexity, and runtime performance. Results show that the SciSpacy pipeline provides significantly faster processing and greater stability when handling large collections of clinical notes, while transformer-based models offer more flexible contextual representations at the cost of increased computational overhead. These findings highlight the trade-offs between lightweight biomedical NLP pipelines and transformer-based models for large-scale clinical text processing.

Keywords: Biomedical NLP, Named Entity Recognition, Clinical Text Processing, SciSpacy, BioBERT

1 Introduction

In this project, Named Entity Recognition (NER) was applied to clinical notes in order to extract biomedical entities that will later be used to construct a clinical knowledge graph. Three different approaches were explored and compared: SciSpacy, BioBERT/PubMedBERT based models, and transformer-based chemical detection models. The experiments were conducted on approximately 5,000 clinical notes collected from publicly available medical transcription samples. These notes contain unstructured clinical narratives describing patient diagnoses, symptoms, and treatments, making them suitable for evaluating biomedical NER pipelines. The goal of this study is to evaluate the practical trade-offs between lightweight biomedical NLP pipelines and transformer-based models when applied to large collections of clinical notes.

2 SciSpacy Pipeline

The first approach used the SciSpacy model `en_ner_bc5cdr_md`. SciSpacy is a biomedical extension of the spaCy natural language processing library designed for scientific and medical text processing (Neumann *et al.*, 2019).

This model is trained on the BC5CDR biomedical dataset and is capable of detecting two important entity types: diseases and chemicals.

The SciSpacy pipeline was straightforward to implement. Each clinical note was processed using the spaCy pipeline, and the detected entities were extracted directly from the model output. The extracted entities were saved into a JSONL file that includes the entity text, label, and character positions in the original note.

During testing, SciSpacy performed reliably and was able to process long clinical notes without any input length limitations. The entity spans were generally clean and required minimal post-processing.

3 Transformer-Based NER Pipeline

To compare with SciSpacy, a transformer-based NER pipeline was implemented using models available through the HuggingFace Transformers library (Wolf *et al.*, 2020). Transformer models are based on the BERT architecture, which is widely used for natural language processing tasks (Devlin *et al.*, 2019).

Two models were used in the transformer pipeline:

- A PubMedBERT-based model for detecting disease entities
- A chemical detection model for identifying medications and other chemical entities

PubMedBERT is a biomedical language model trained entirely on biomedical literature from PubMed (Gu *et al.*, 2021). Similar biomedical transformer models such as BioBERT have shown strong performance in biomedical text mining tasks (Lee *et al.*, 2020).

Since these models were trained on different datasets, a single transformer model was not able to detect both diseases and chemicals reliably. Therefore, the outputs of the two models were combined to produce the final entity list.

The transformer models were applied to the same dataset of clinical notes, and the extracted entities were saved in the same JSONL format as the SciSpacy pipeline.

4 Implementation Challenges

Several practical challenges were encountered while implementing the transformer-based pipeline.

First, transformer models have a maximum input size of approximately 512 tokens due to the underlying BERT architecture (Devlin *et al.*, 2019). Some clinical notes were longer than this limit, which caused runtime errors. To address this issue, a preprocessing function was implemented to truncate very long notes before sending them to the model.

Second, transformer tokenization sometimes splits words into subword tokens. For example, a medication name such as “lisinopril” may be split into multiple pieces. A custom post-processing function was implemented to merge these tokens back into a single word.

Third, some transformer models produced partial entity spans or false positive predictions. For example, the phrase “type 2 diabetes” was sometimes partially detected as “type” only.

5 Experimental Environment

All experiments were conducted on a system equipped with an NVIDIA RTX 3070 Ti Mobile GPU, an AMD Ryzen 6900HX CPU and 16GB of RAM. The experiments were implemented using Python with the HuggingFace Transformers library and SciSpacy models.

Although a GPU was available, the majority of experiments were executed on CPU due to compatibility and memory considerations when processing long clinical notes. Runtime

performance therefore reflects a CPU-based execution environment.

6 Performance Comparison

During runtime evaluation, the transformer-based pipeline was significantly slower than the SciSpacy pipeline. Processing approximately 5000 clinical notes required more than eight hours when executed on CPU, with an average processing time of approximately 6 seconds per note.

In contrast, the SciSpacy pipeline processed the dataset much faster and required fewer additional preprocessing steps.

6.1 Runtime Statistics

To evaluate the computational cost of each NER approach, runtime measurements were collected during experimentation. The total execution time was recorded and divided by the number of processed clinical notes in order to obtain the average runtime per note.

The SciSpacy pipeline demonstrated the fastest performance among the evaluated methods. Processing the dataset required approximately 262.47 seconds in total, corresponding to an average runtime of roughly 0.05 seconds per clinical note.

In comparison, the transformer-based pipelines required substantially more computation time. The BioBERT/PubMedBERT-based pipeline required approximately 6.34 seconds per note on average when processing the dataset. The ClinicalBERT model showed improved efficiency compared to the BioBERT-based pipeline, requiring a total runtime of approximately 1085.43 seconds, corresponding to an average runtime of 0.2186 seconds per clinical note.

Overall, these results highlight the efficiency advantage of the SciSpacy pipeline for large-scale clinical note processing. While transformer-based models provide stronger contextual language representations, their higher computational cost introduces significant overhead when processing large datasets.

6.2 NER Model Comparison

Table 1: Comparison of NER models used in this project

Aspect	SciSpacy	BioBERT / Pub-MedBERT	ClinicalBERT
Model Type	Traditional biomedical NLP pipeline	Transformer (BERT-based)	Transformer (clinical-domain BERT)
Setup Complexity	Very simple pipeline setup	Moderate complexity	Moderate complexity
Entity Types	Disease and chemical entities	Disease and chemical using separate models	General clinical entities
Handling Long Notes	No input length limitation	Limited to ~512 tokens	Limited to ~512 tokens but pipeline handled it automatically
Tokenization Behaviour	Clean entity spans	Subword token splitting	Subword token splitting
Post-processing Required	Minimal	Token merging required	Token merging sometimes required
Runtime Performance	Fast	Slow on CPU	Slow on CPU
Engineering Complexity	Low	High	High
Stability in Clinical Notes	High	Moderate	Moderate
Overall Suitability	Efficient for large datasets	Flexible but computationally expensive	Useful for clinical language context

7 Summary of Findings

Overall, the results indicate that SciSpacy provides a more stable and efficient solution for large-scale biomedical entity extraction from clinical notes. The transformer-based pipeline offers flexibility and the potential for deeper contextual understanding, but it requires additional engineering steps and greater computational resources.

For this project, SciSpacy was therefore used as the primary NER pipeline, while the transformer-based approach was retained as an experimental comparison.

References

- [1] Neumann, M., King, D., Beltagy, I. and Ammar, W., 2019. ScispaCy: Fast and Robust Models for Biomedical Natural Language Processing. Proceedings of the 18th BioNLP Workshop and Shared Task. Florence, Italy, pp.319–327. Available at: <https://doi.org/10.18653/v1/W19-5034>
- [2] Devlin, J., Chang, M-W., Lee, K. and Toutanova, K., 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT), pp.4171–4186. Available at: <https://doi.org/10.18653/v1/N19-1423>
- [3] Lee, J., Yoon, W., Kim, S., Kim, D., Kim, S., So, C.H. and Kang, J., 2020. BioBERT: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4), pp.1234–1240. Available at: <https://doi.org/10.1093/bioinformatics/btz682>
- [4] Gu, Y., Tinn, R., Cheng, H., Lucas, M., Usuyama, N., Liu, X., Naumann, T., Gao, J. and Poon, H., 2021. Domain-Specific Language Model Pretraining for Biomedical Natural Language Processing. *ACM Transactions on Computing for Healthcare*, 3(1), pp.1–23. Available at: <https://doi.org/10.1145/3458754>
- [5] Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M. and Brew, J., 2020. Transformers: State-of-the-Art Natural Language Processing. Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations, pp.38–45. Available at: <https://doi.org/10.18653/v1/2020.emnlp-demos.6>